

Kubernetes Essentials

Rey Lejano @ DevWeek 2024



kubernetes

Who am I?



- Solutions Architect @ Red Hat
- Kubernetes SIG Docs co-chair
- Kubernetes SIG Security subproject lead
- Kubernetes v1.23 Release Lead
- CNCF Ambassador
- DevOps Institute Ambassador

@reylejano





Agenda

- Container basics
- Kubernetes history
- Kubernetes architecture
- Kubernetes resources to deploy an application and get outside traffic to the application
- Tutorial

<https://github.com/reylejano/kubernetes-essentials-devweek>

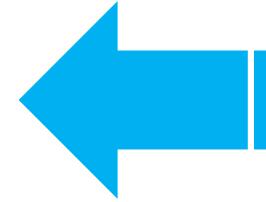
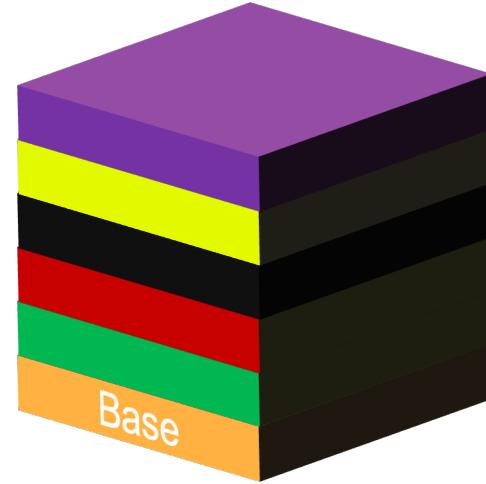
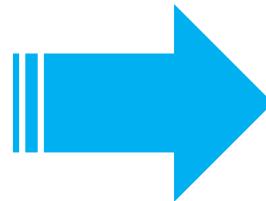
Container Basics



Containerization



```
1 package main
2
3 import (
4     "fmt"
5     "net/http"
6 )
7
8 func hello(w http.ResponseWriter, req *http.Request) {
9     fmt.Fprintf(w, "Hello, world!\n")
10 }
11
12 func headers(w http.ResponseWriter, req *http.Request) {
13     for name, headers := range req.Header {
14         for _, h := range headers {
15             fmt.Fprintf(w, "%v: %v\n", name, h)
16         }
17     }
18 }
19
20 func main() {
21     http.HandleFunc("/hello", hello)
22     http.HandleFunc("/headers", headers)
23     http.ListenAndServe(":80", nil)
24 }
```

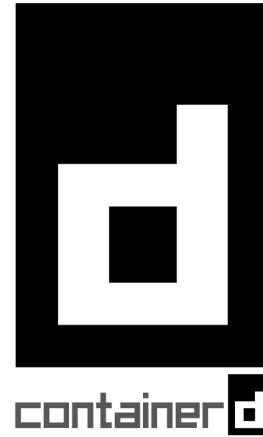
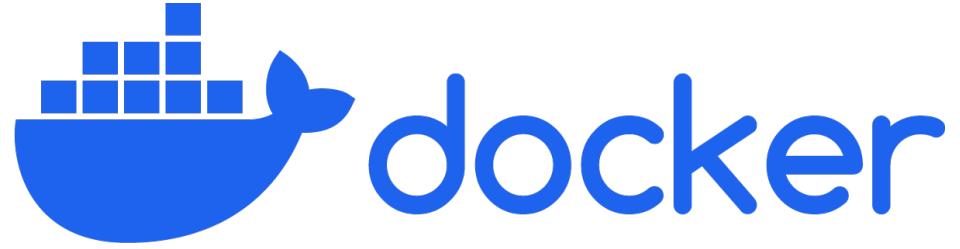


Containers are the implementation of
cgroups & namespaces

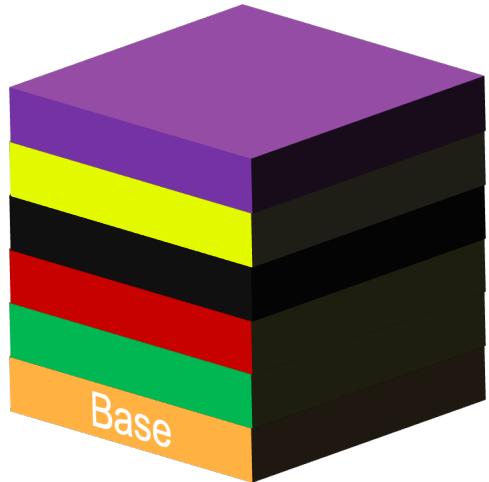


kubernetes

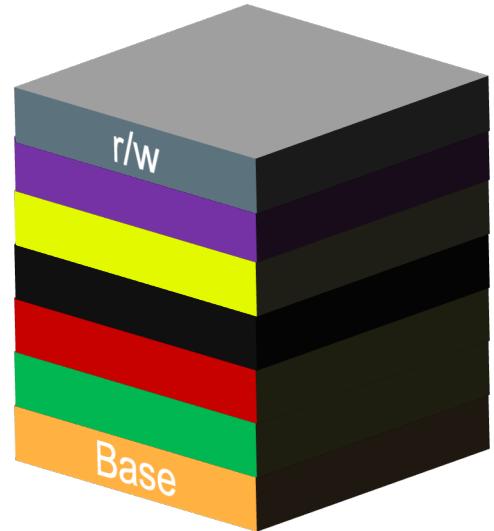
Container Runtime



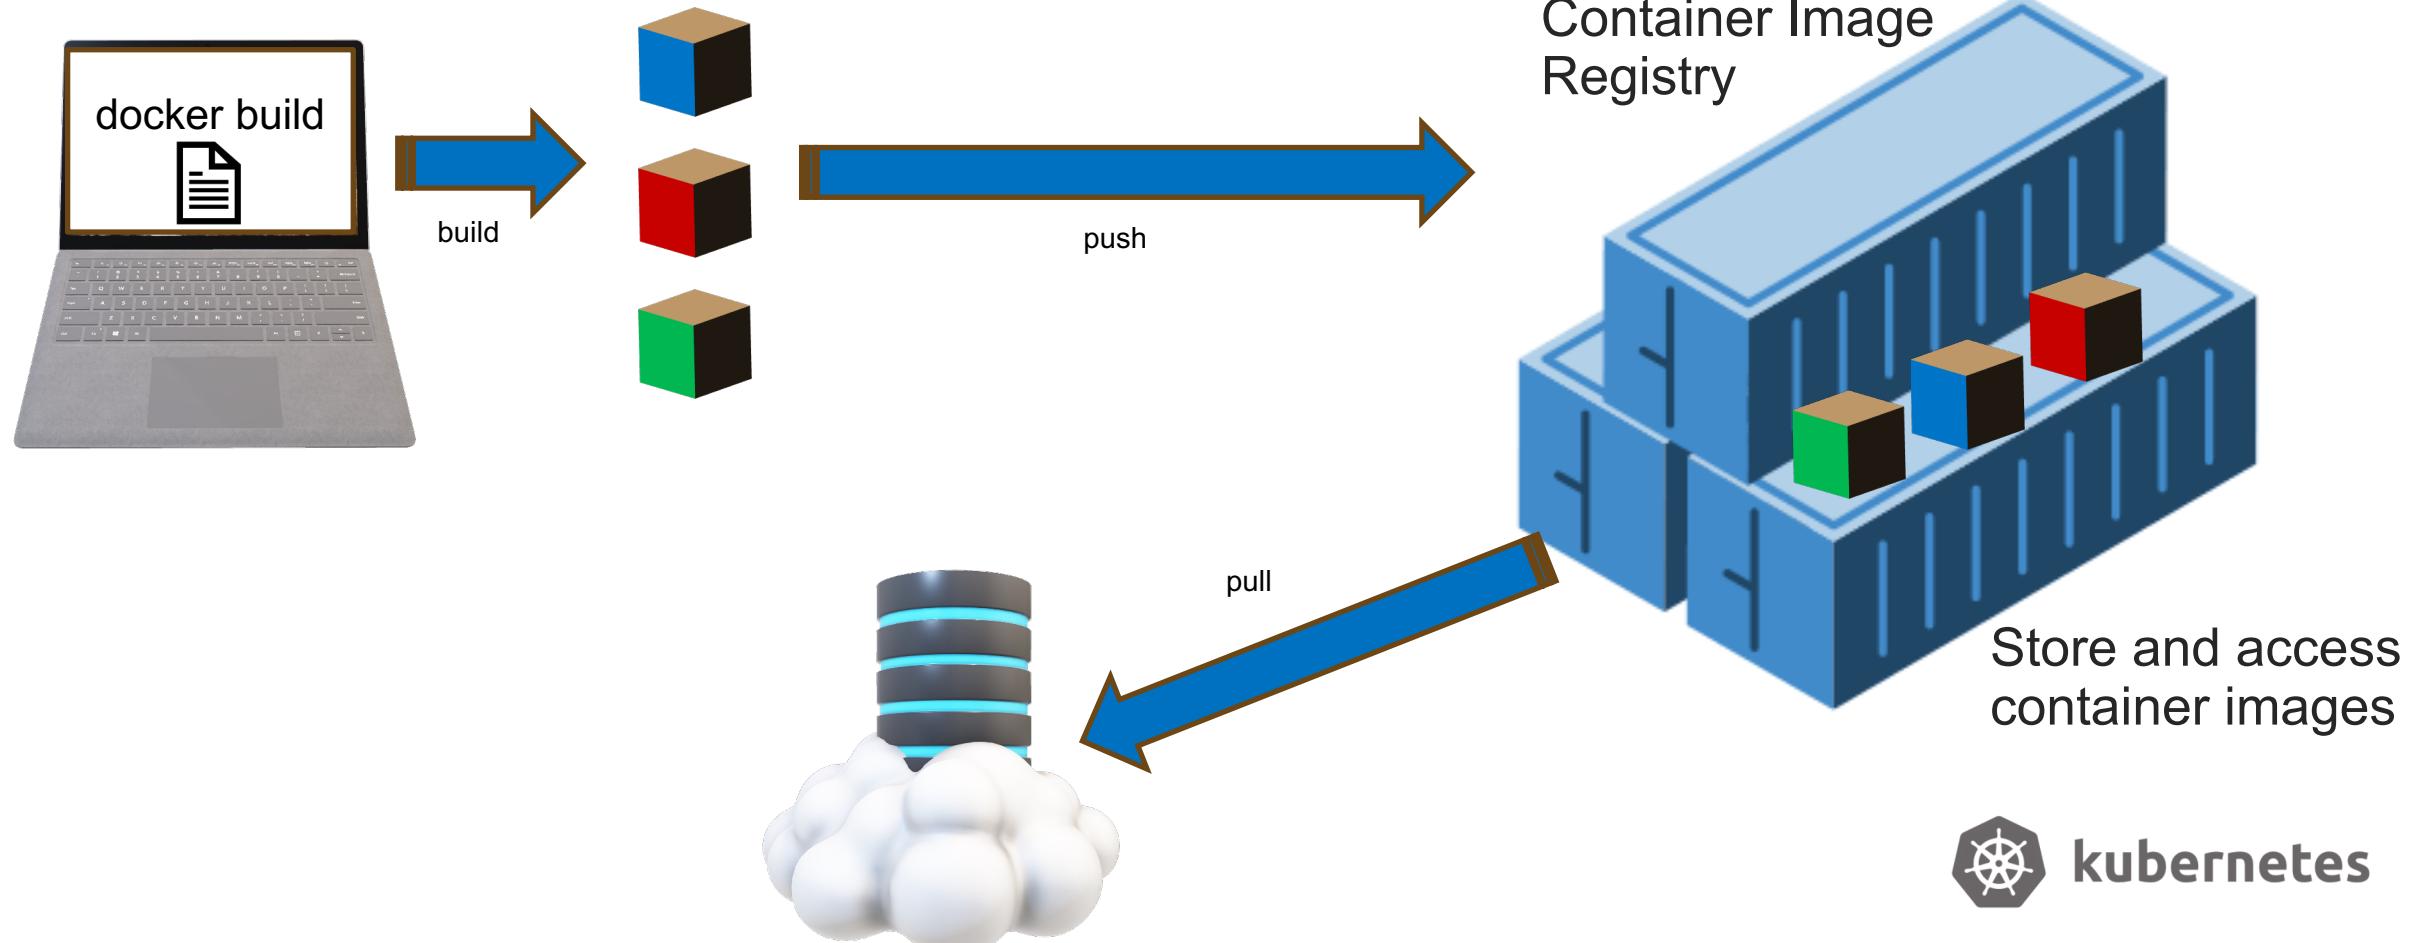
runc



Container runtime runs executes containers
based on container images



Container Registry



Cloud Native

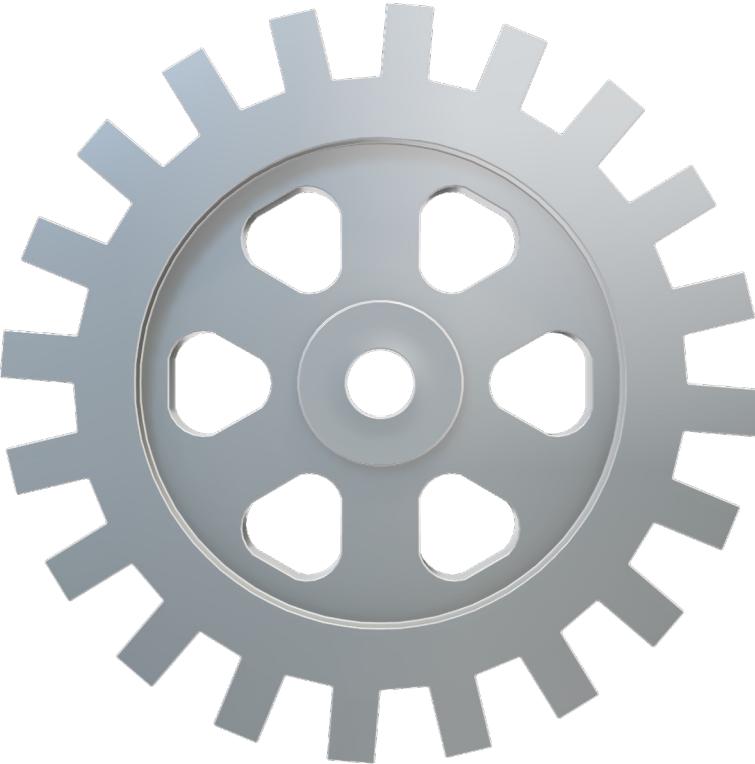
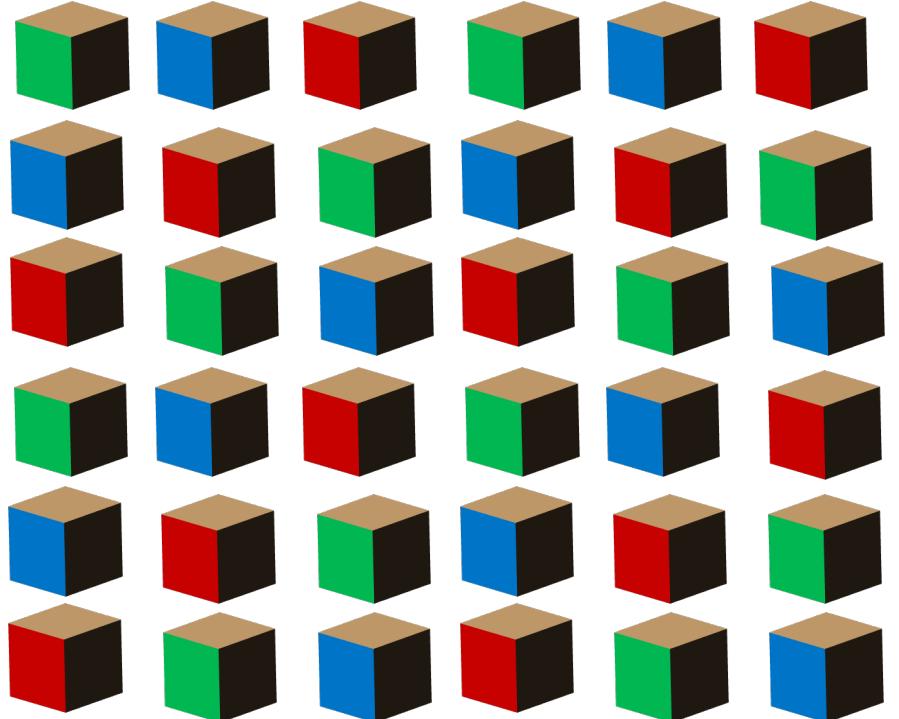


- Containers are foundation of cloud native applications – take advantage of cloud computing: runs anywhere, takes advantage of distributed compute and scales
- Decouples apps to specific environments / microservices / [12 factor app](#)
- Open source
- Runs anywhere
- Interoperates with other tools and common conventions

Container Orchestration



Container Orchestration

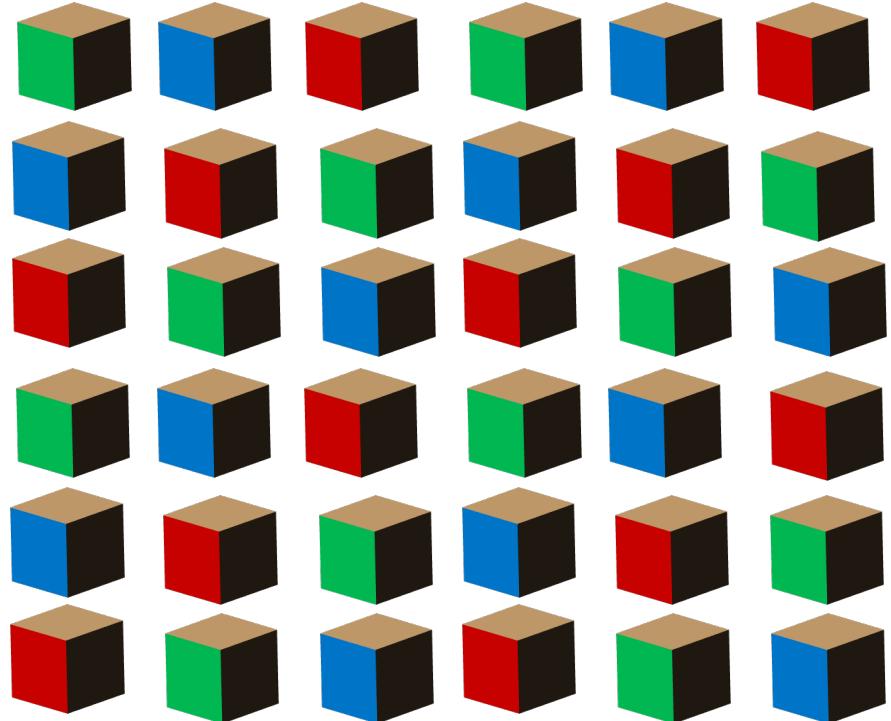


How do we solve...
- Automation?
- Orchestration?
- Security?
- Networking?



kubernetes

Container Orchestration



kubernetes





Kubernetes

- Container orchestrator that deploys, manages, scales containerized workloads on distributed compute at-scale.
- Greek for Helmsman or Pilot **κυβερνήτης**
- Utilizes the kernel of distributed systems so it can run "anywhere" AKA "Linux kernel" of distributed systems
- Abstracts the underlying hardware of distributed compute and use the same API on every cloud and on-prem to manage workloads
- Works so that the actual state matches the desired state (self-healing)

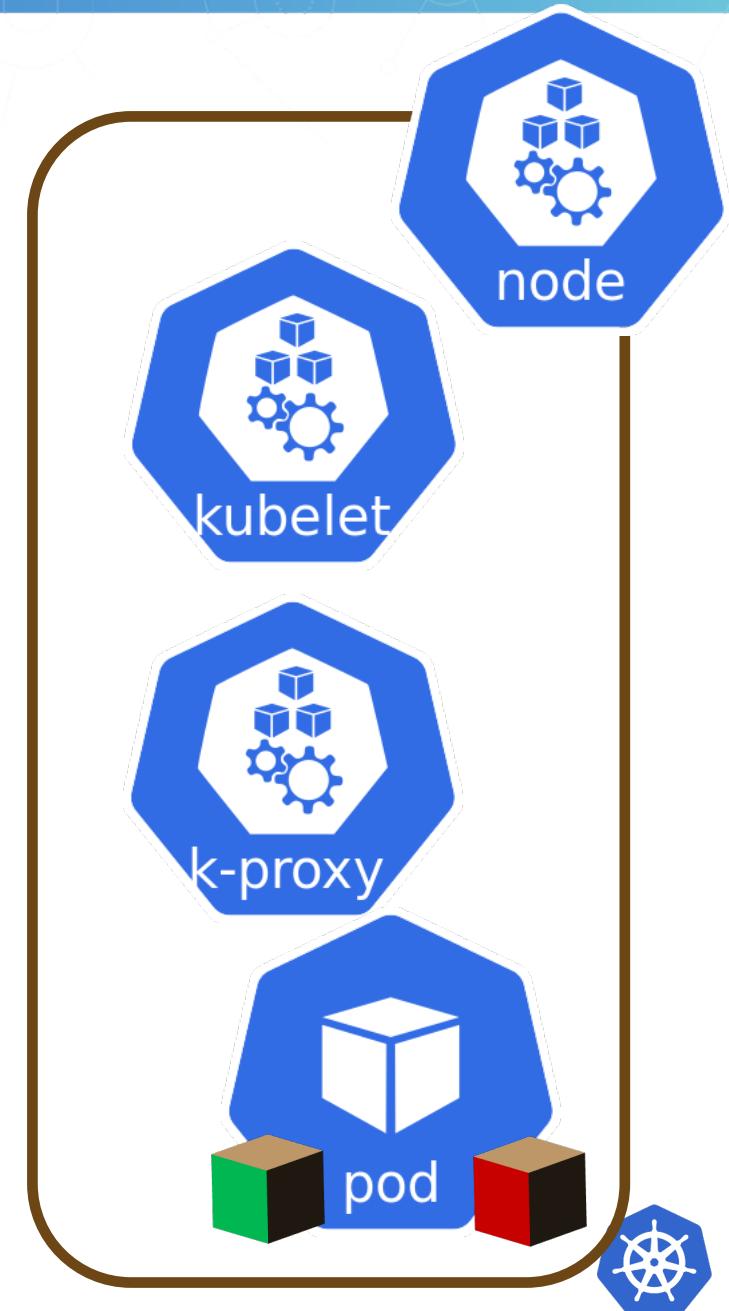
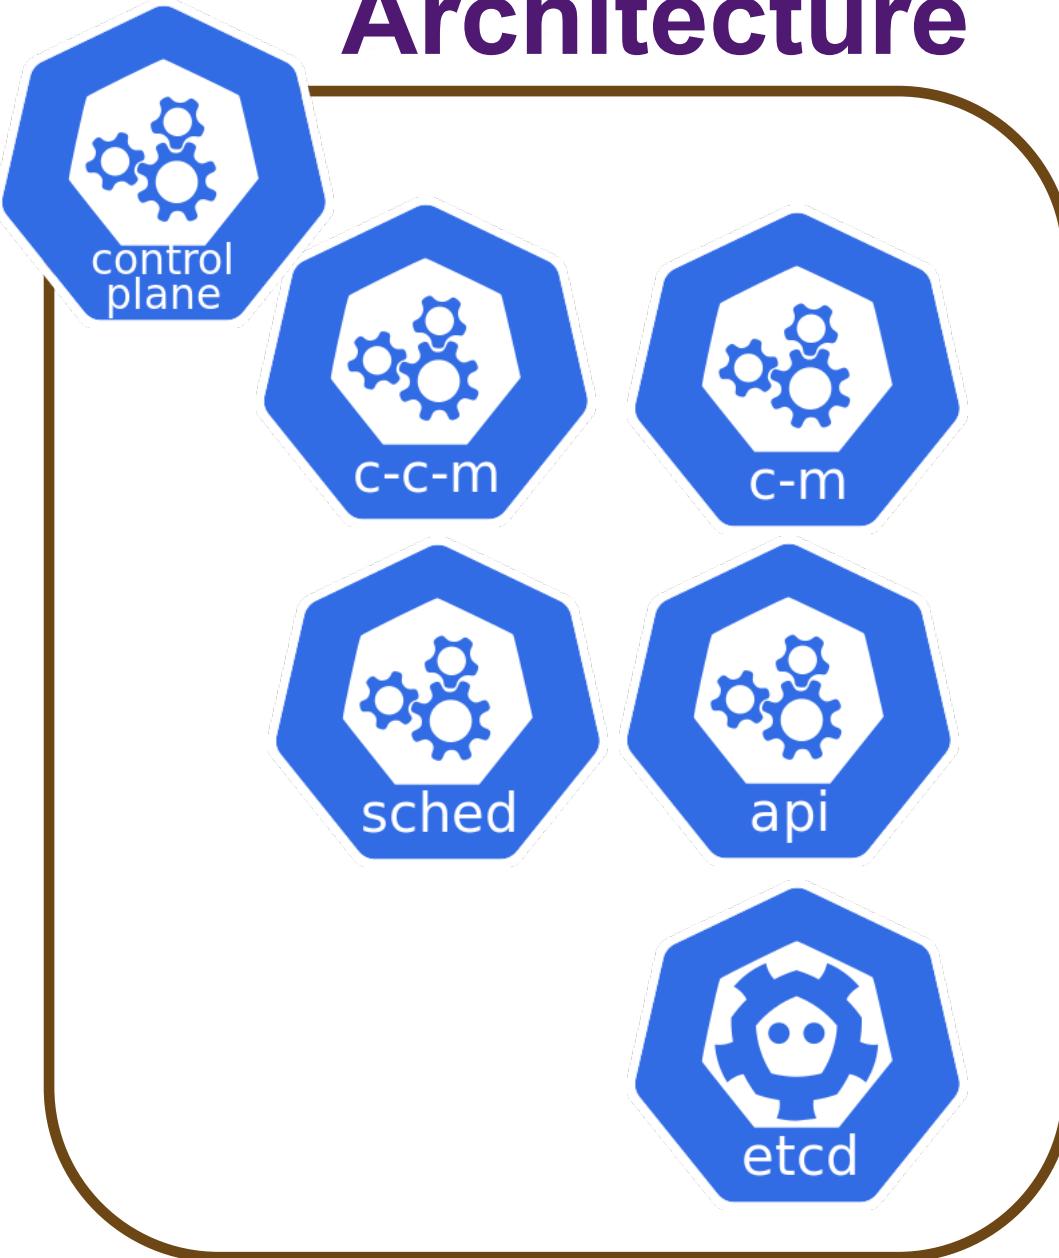




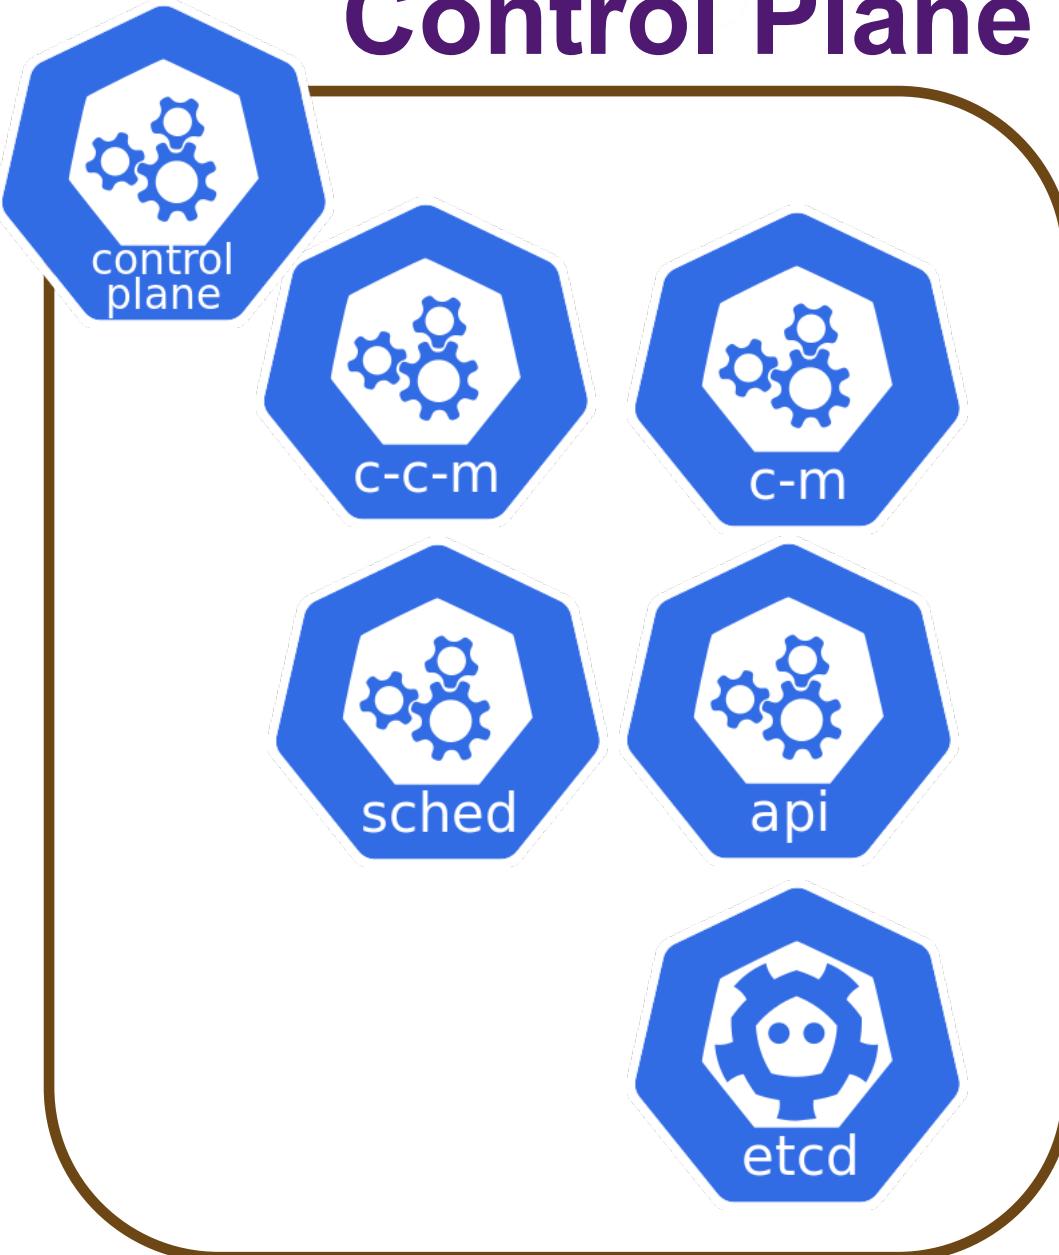
Kubernetes History

- Public release June 2014
- Created by Google built from lessons learned from Google's Borg & Omega container orchestrators. White paper titled [Borg, Omega, and Kubernetes](#)
- June 2015 Open Container Initiative (OCI) - standardization
- July 2015 Kubernetes v1.0. Google and the Linux Foundation helped form the Cloud Native Computing Foundation (CNCF) and donated Kubernetes to the CNCF as the seed project
- 2016 helm, kubeadm released. Notable case study on [Pokémon Go on Kubernetes](#)
- 2017 Notable talk at KubeCon + CloudNativeCon Europe Keynote from Open AI titled [Building the Infrastructure that Powers the Future of AI](#)
- 10 year birthday on June 6, 2024

Architecture



Control Plane

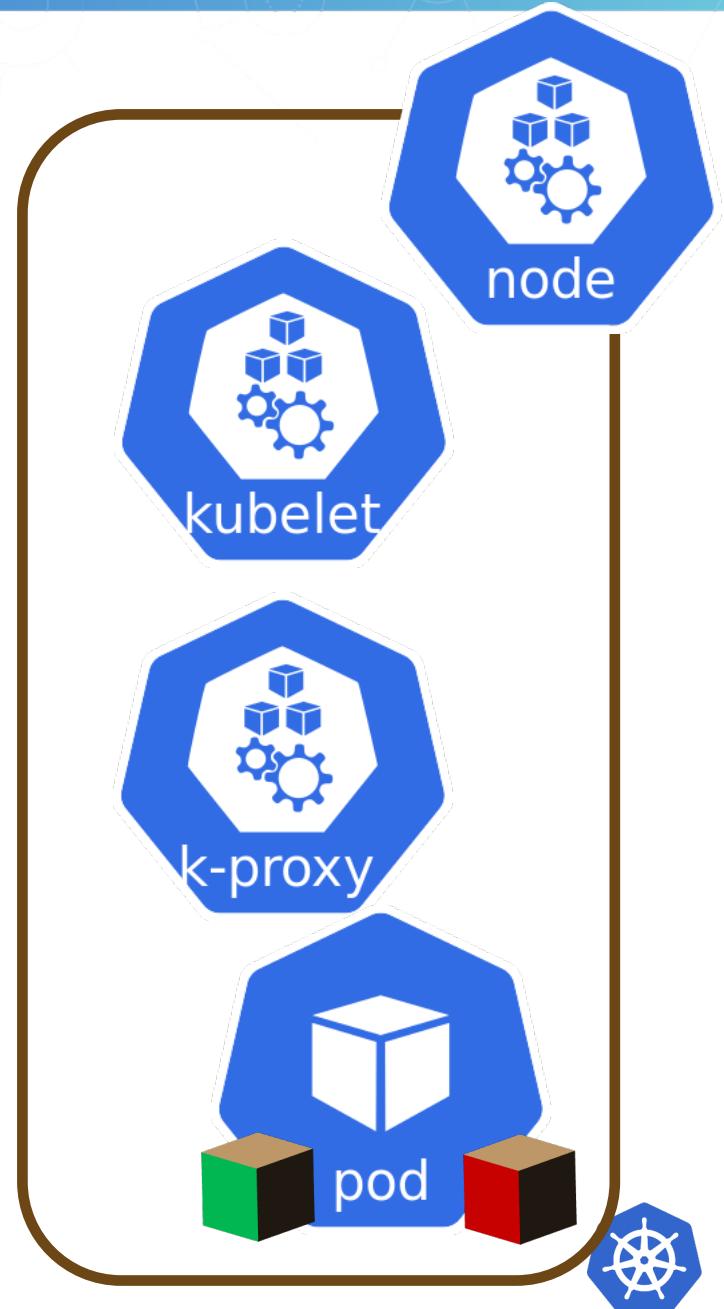


- kube-apiserver – REST interface, handles all interaction between users, nodes, etcd, 3rd party
- etcd – cluster data store (key-value)
- cloud controller manager – provides cloud-provider specific integration capability
- controller manager – daemon that manages all component control loops, monitors cluster state and steers the cluster towards desired state
- scheduler – binds (schedules) pods to nodes

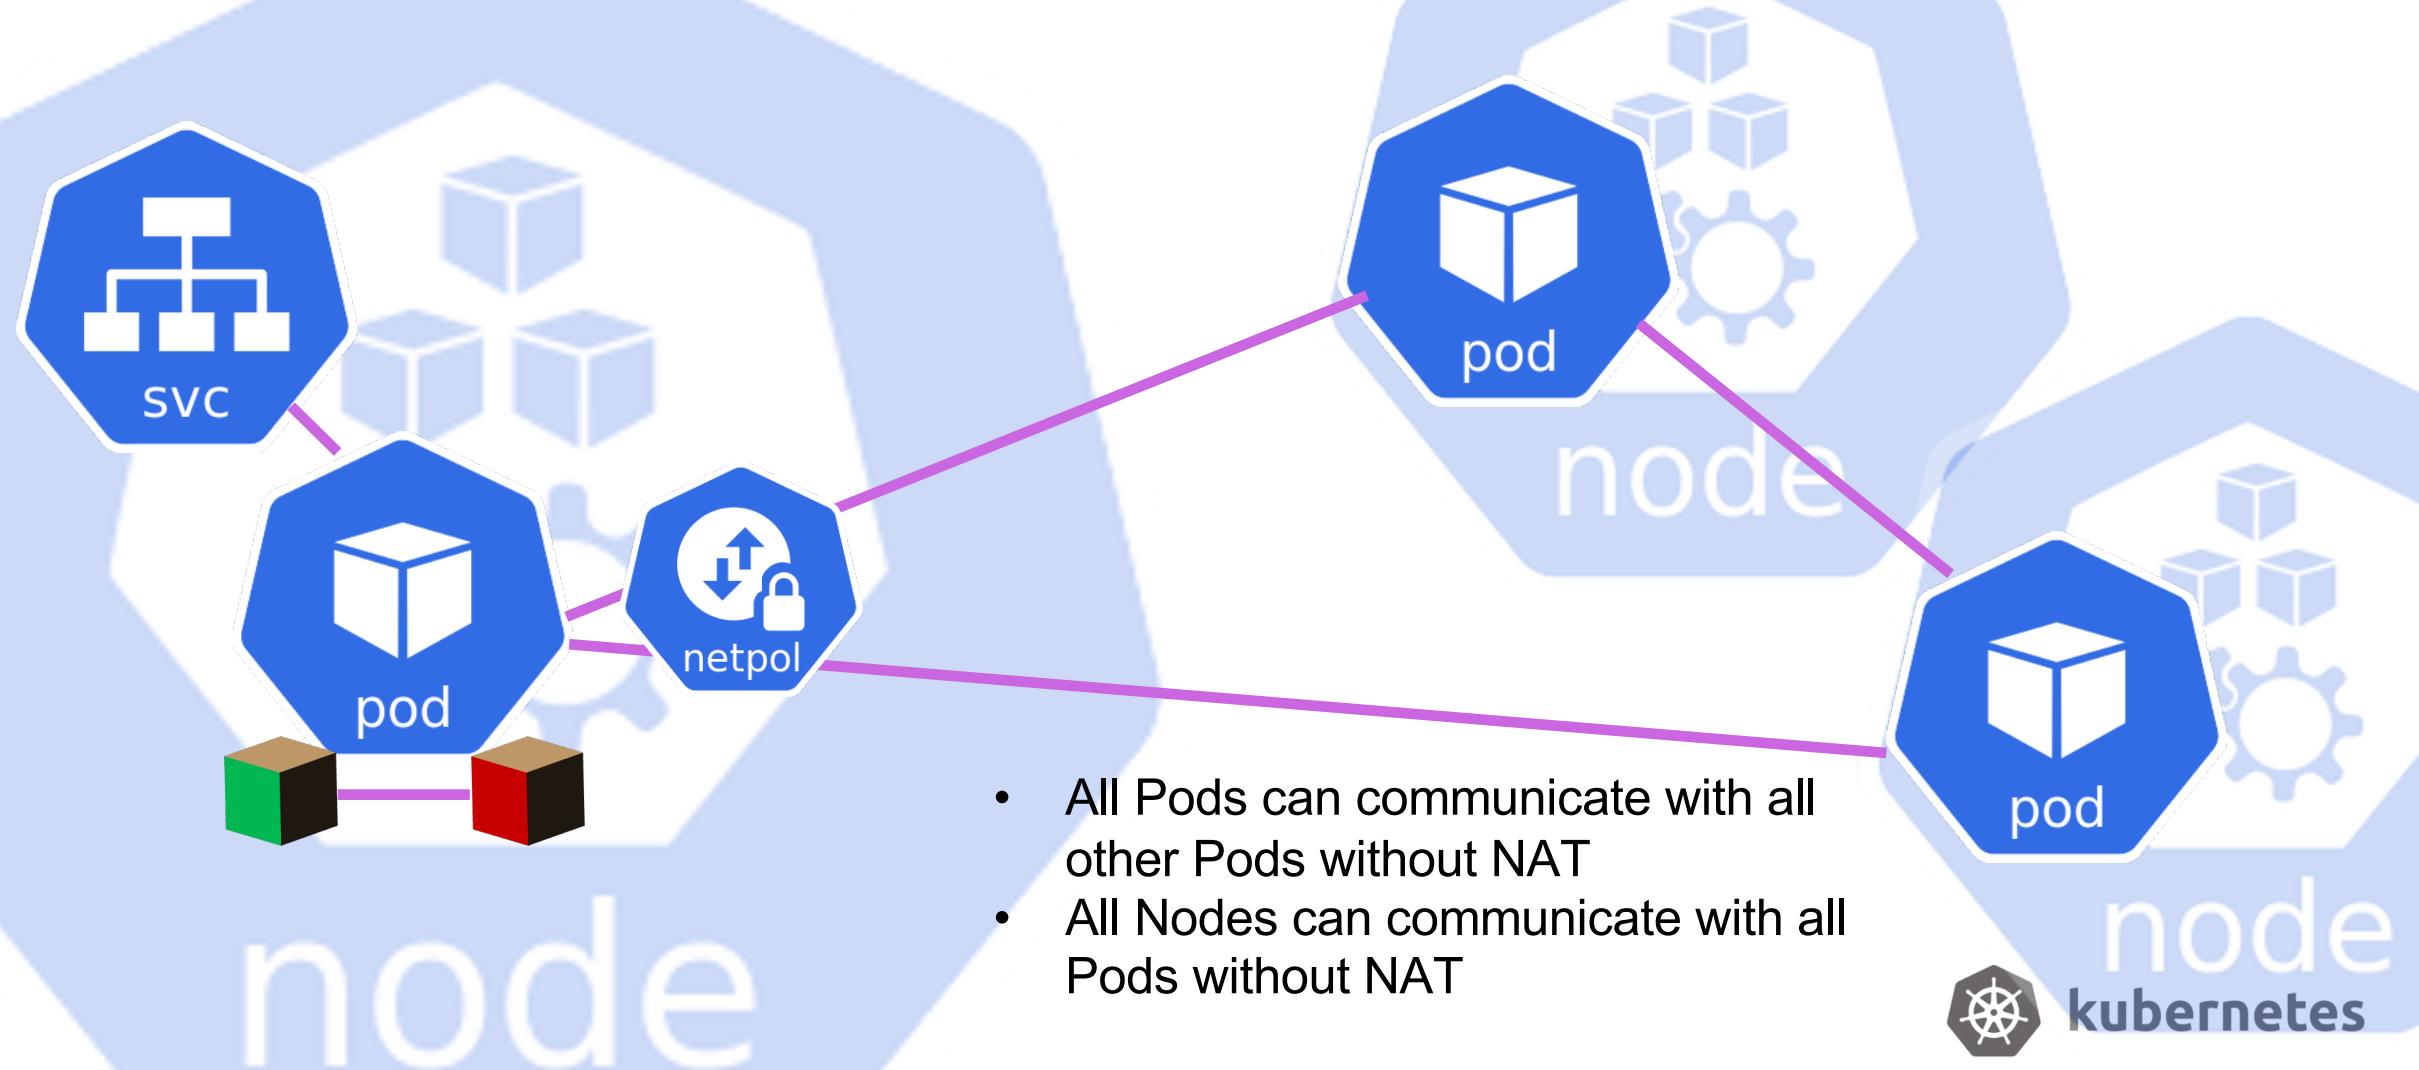


Worker Node

- kubelet – node agent, responsible for Pod lifecycle
- kube-proxy – manages network rules on each node and performs forwarding and load balancing for Services
- container runtime



Kubernetes Networking



Kubernetes Resources



Deploying on Kubernetes

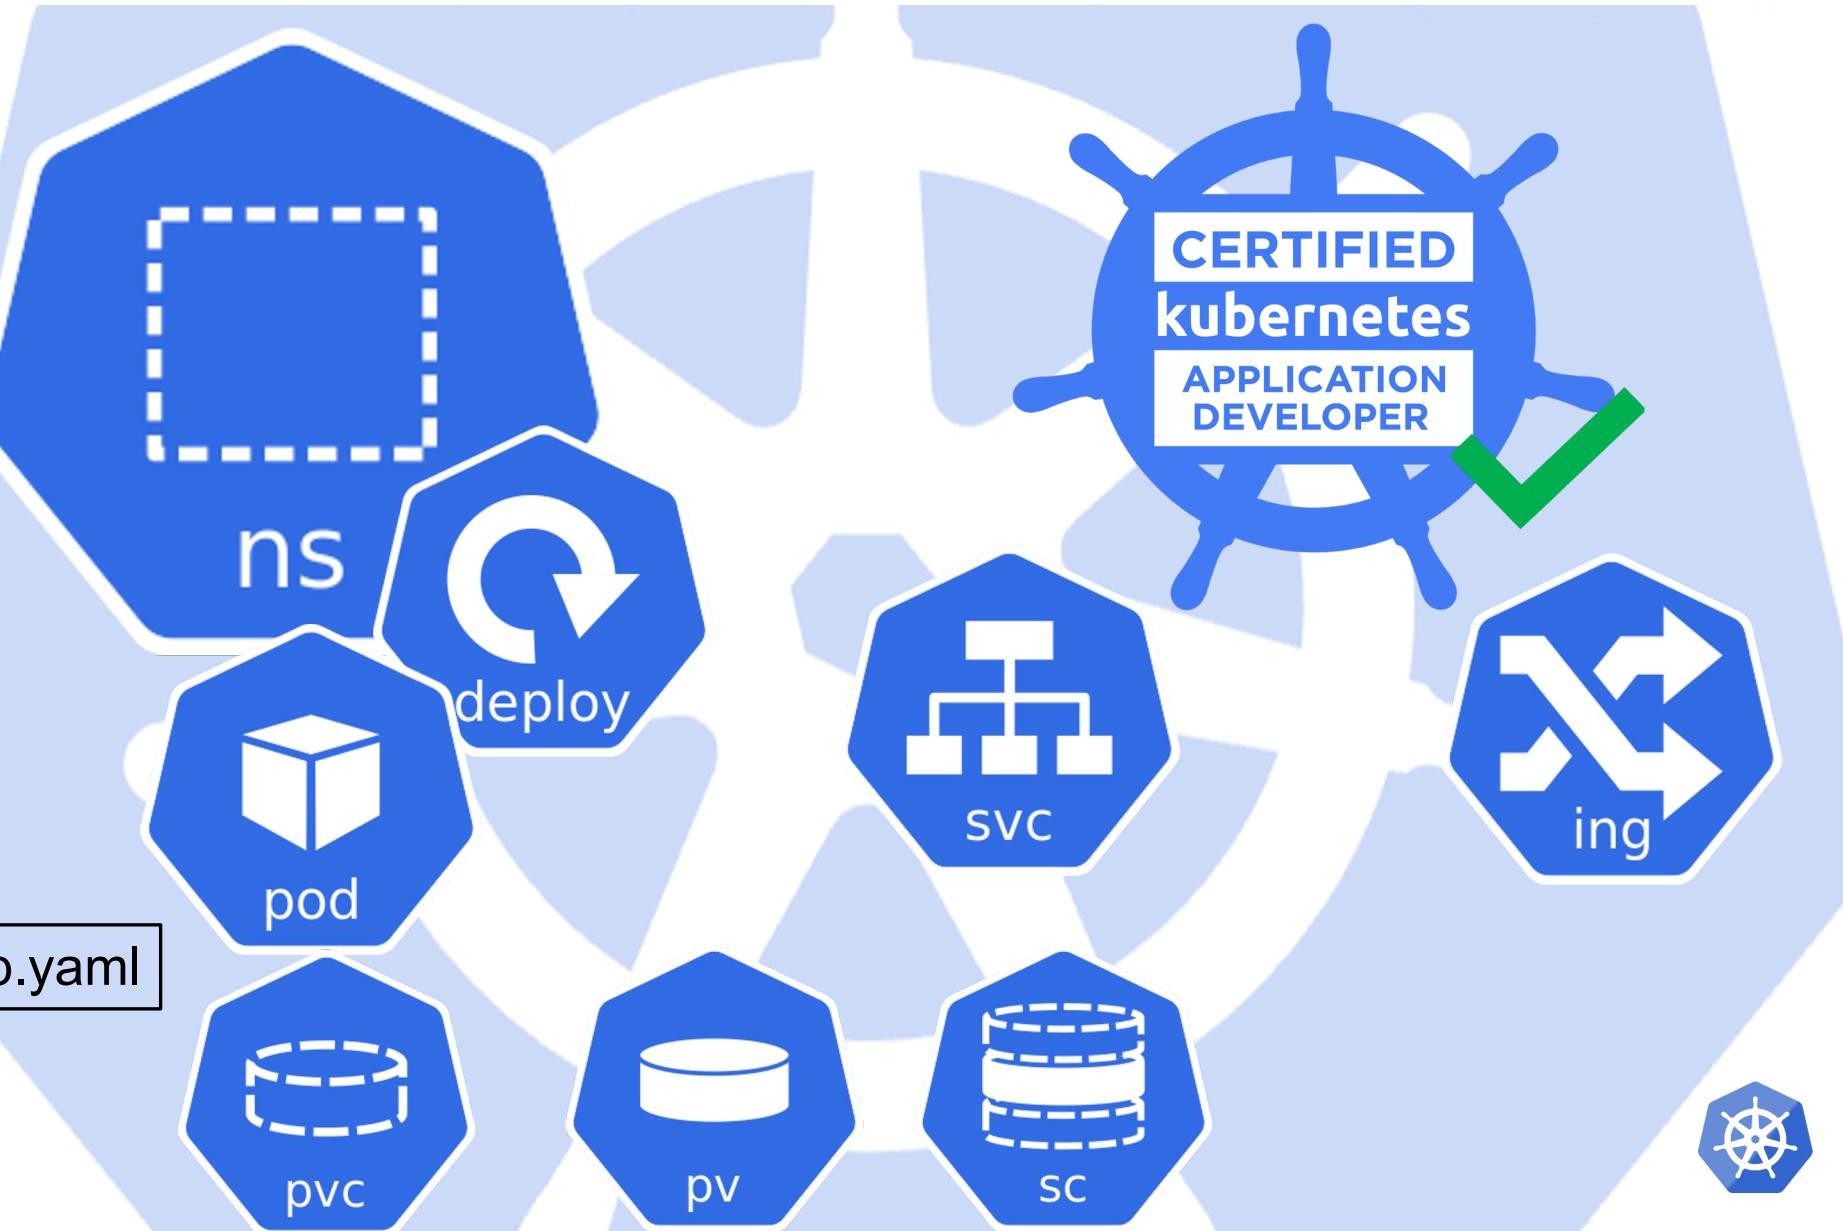
```
1 package main
2
3 import (
4     "fmt"
5     "net/http"
6     "os"
7     "log"
8 )
9
10 func main() {
11     logger := log.New(os.Stdout, "http:", log.LstdFlags)
12     logger.Println("Server is starting...")
13     http.HandleFunc("/", handler)
14     http.ListenAndServe(":8080", nil)
15 }
16
17 func handler(w http.ResponseWriter, r *http.Request) {
18     hostname, err := os.Hostname()
19     if err != nil {
20         fmt.Println(err)
21         os.Exit(1)
22     }
23     fmt.Fprintf(w, "Hostname is "+hostname)
```



container image



app.yaml



Namespace



```
apiVersion: v1
kind: Namespace
metadata:
  name: mynamespace
```

- Provides isolation of resources
- Divide cluster resources among users, applications, projects, teams
- Can have limits and quotas
- Well-known namespaces:
 - default
 - kube-system – created for the Kubernetes system



- Divide cluster resources with resource quotas
- Limit number of resources e.g. # of Pods, Deployments, Services, etc
- CPU & Memory requests and limits

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: myquota
  namespace: mynamespace
spec:
  hard:
    memory: 2Gi
    pods: "20"
```



kubernetes

Pod



```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace:
    mynamespace
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.25.4
      ports:
        - containerPort: 80
```

- Smallest deployable unit
- Composed of 1 or more containers with shared resources like networking, storage
- Pods get a virtual IP (10.42.0.0/16) and a DNS record with the following format: <pod-IPv4>.<namespace>.pod.cluster.local
- Containers of a Pod are co-located

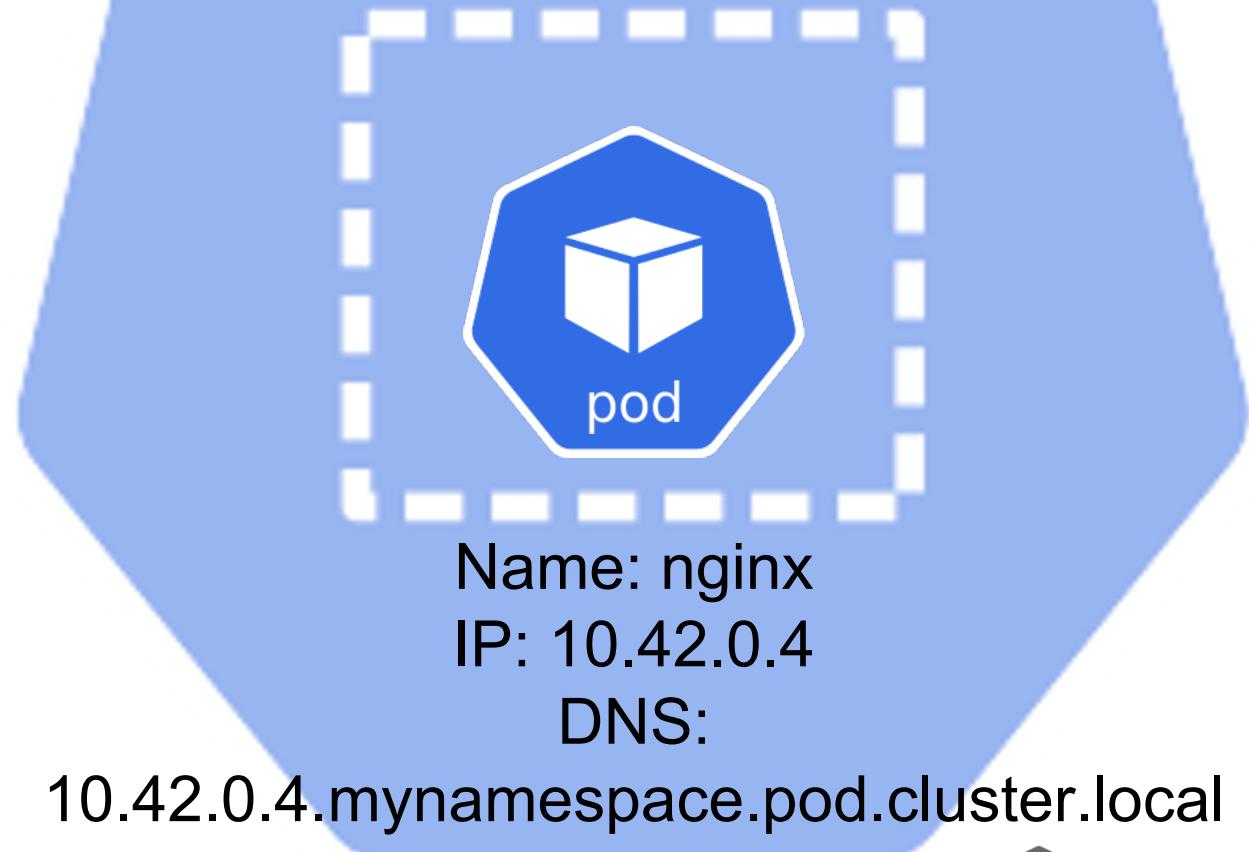


kubernetes

Pod



```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace:
    mynamespace
  labels:
    app: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.25.4
    ports:
    - containerPort: 80
```



kubernetes

Pod



```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace:
    mynamespace
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:1.25.4
      ports:
        - containerPort: 80
```



Don't manage
Pods
manually!



Name: nginx
IP: 10.42.0.4
DNS:

10.42.0.4.mynamespace.pod.cluster.local



kubernetes

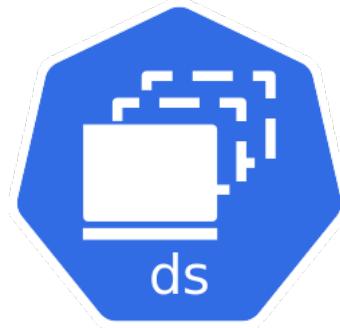
Workload Resources



Deployment



Job



DaemonSet



StatefulSet



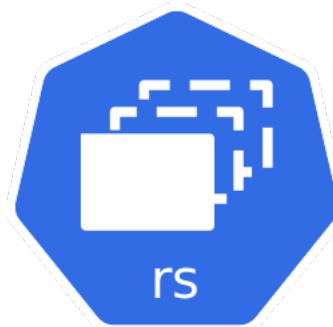
Cronjob



Deployment



- Declarative updates to Pods (via ReplicaSets) like updates, rollouts, rollbacks
- Set the desired state in a Deployment
- Controller changes the actual state to the desired state at a controlled rate
 - (default) no less than 75% and no more than 125% of the desired number of Pods are up



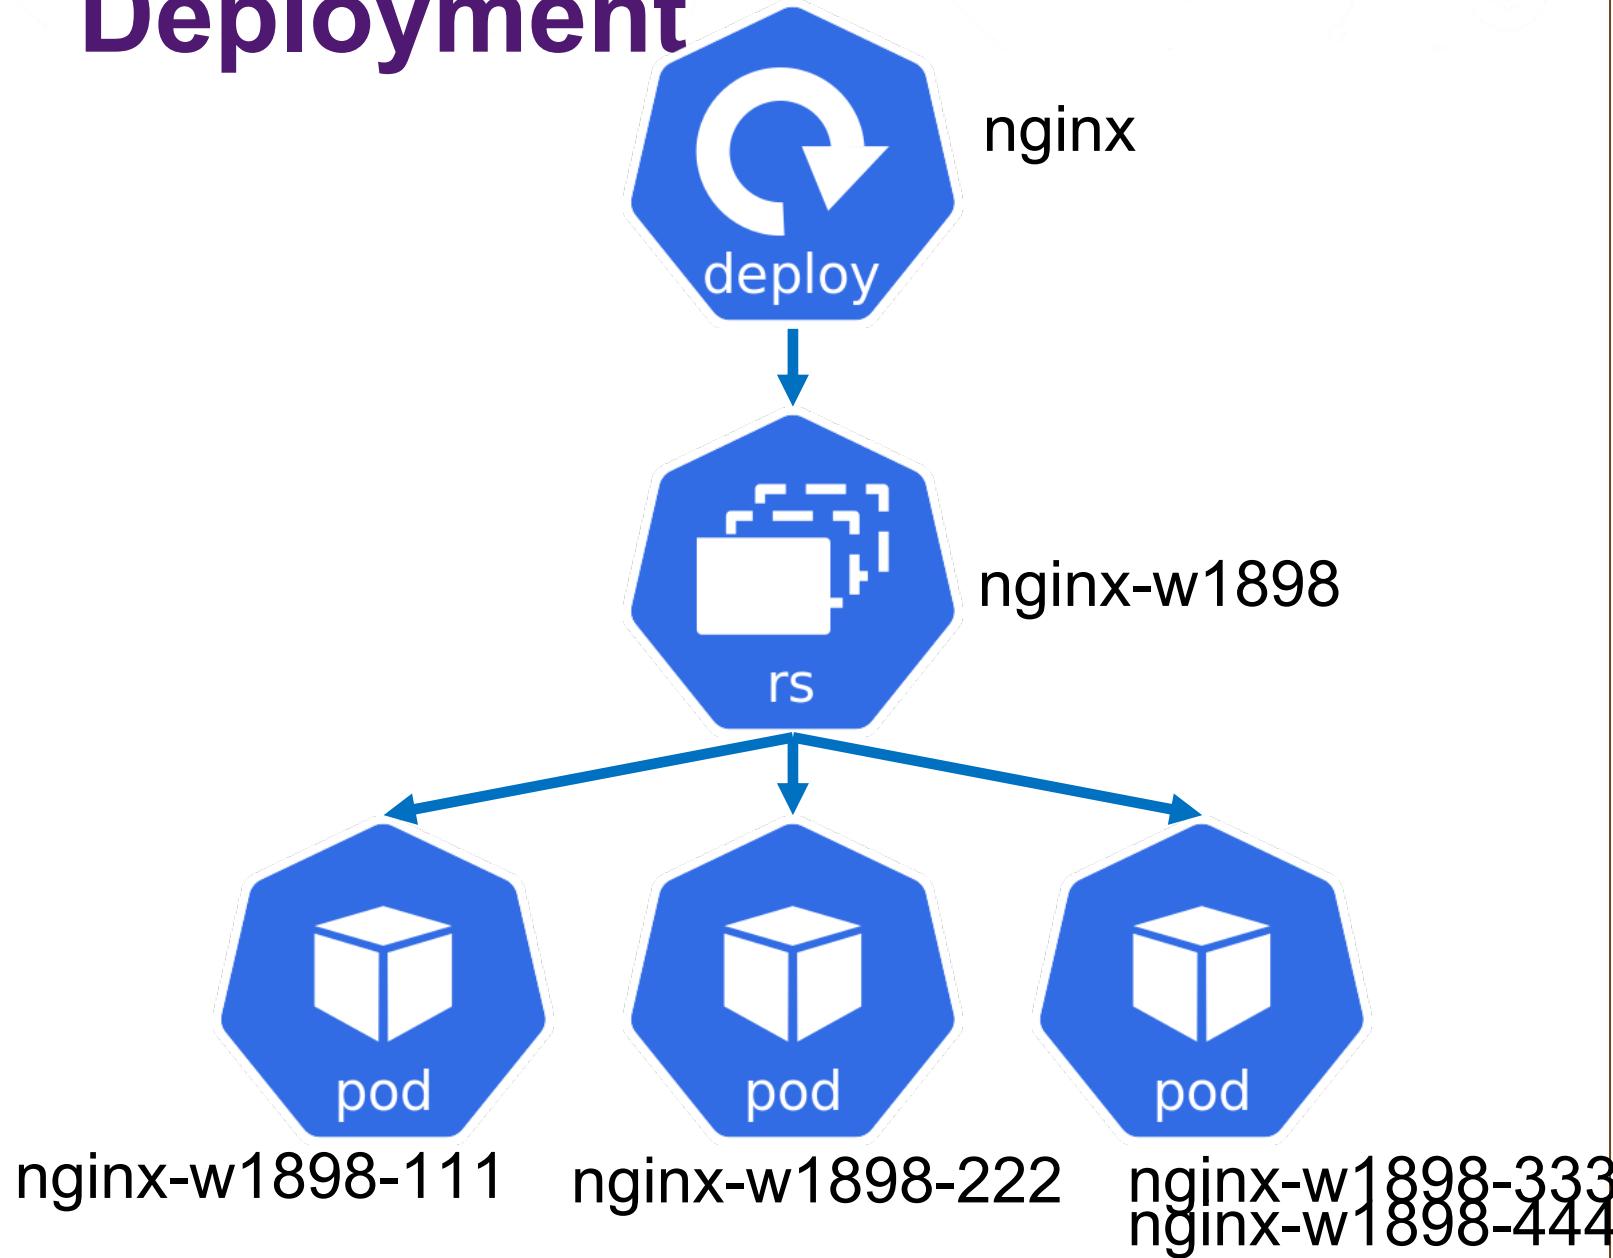
ReplicaSet

Maintains a stable set of Pod replicas

```
apiVersion: v1
kind: Deployment
metadata:
  name: nginx
  namespace: mynamespace
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.25.4
          ports:
            - containerPort: 80
```



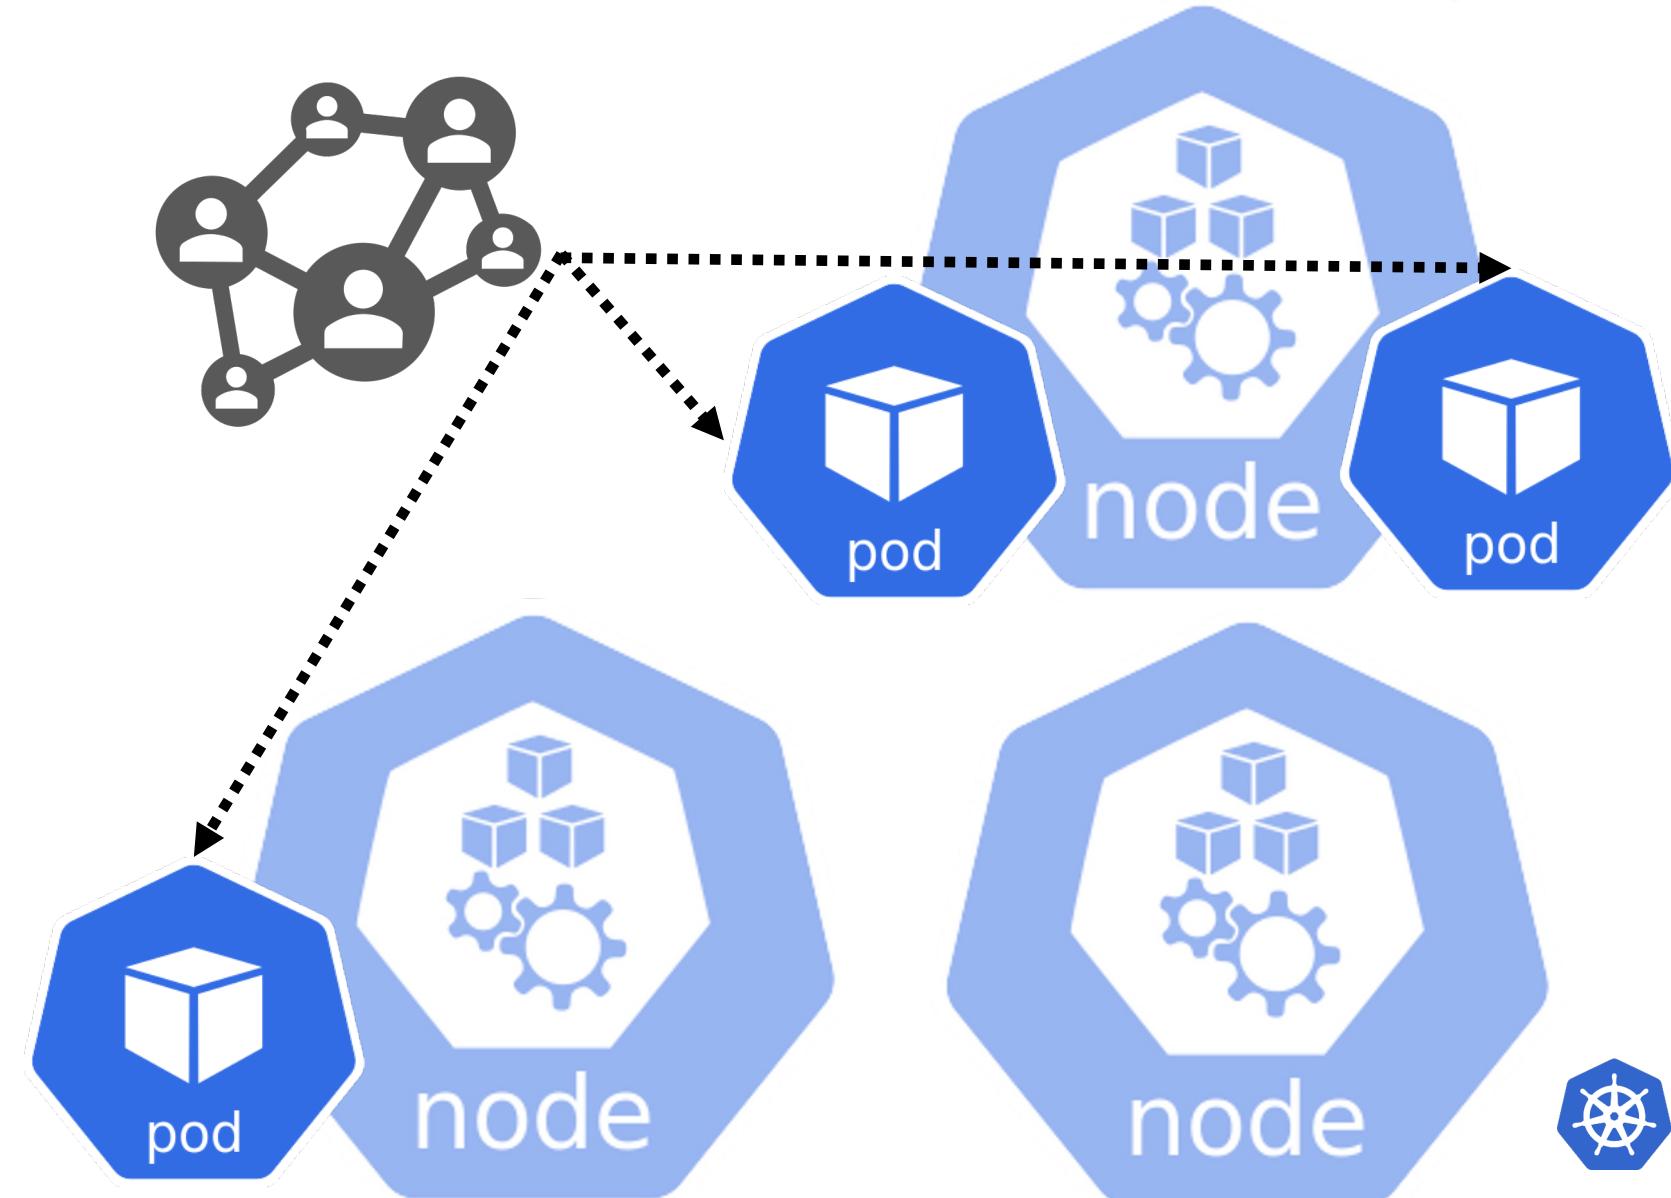
Deployment



```
apiVersion: v1
kind: Deployment
metadata:
  name: nginx
  namespace: mynamespace
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.25.4
          ports:
            - containerPort: 80
```



Getting Traffic to Your Application / Microservice



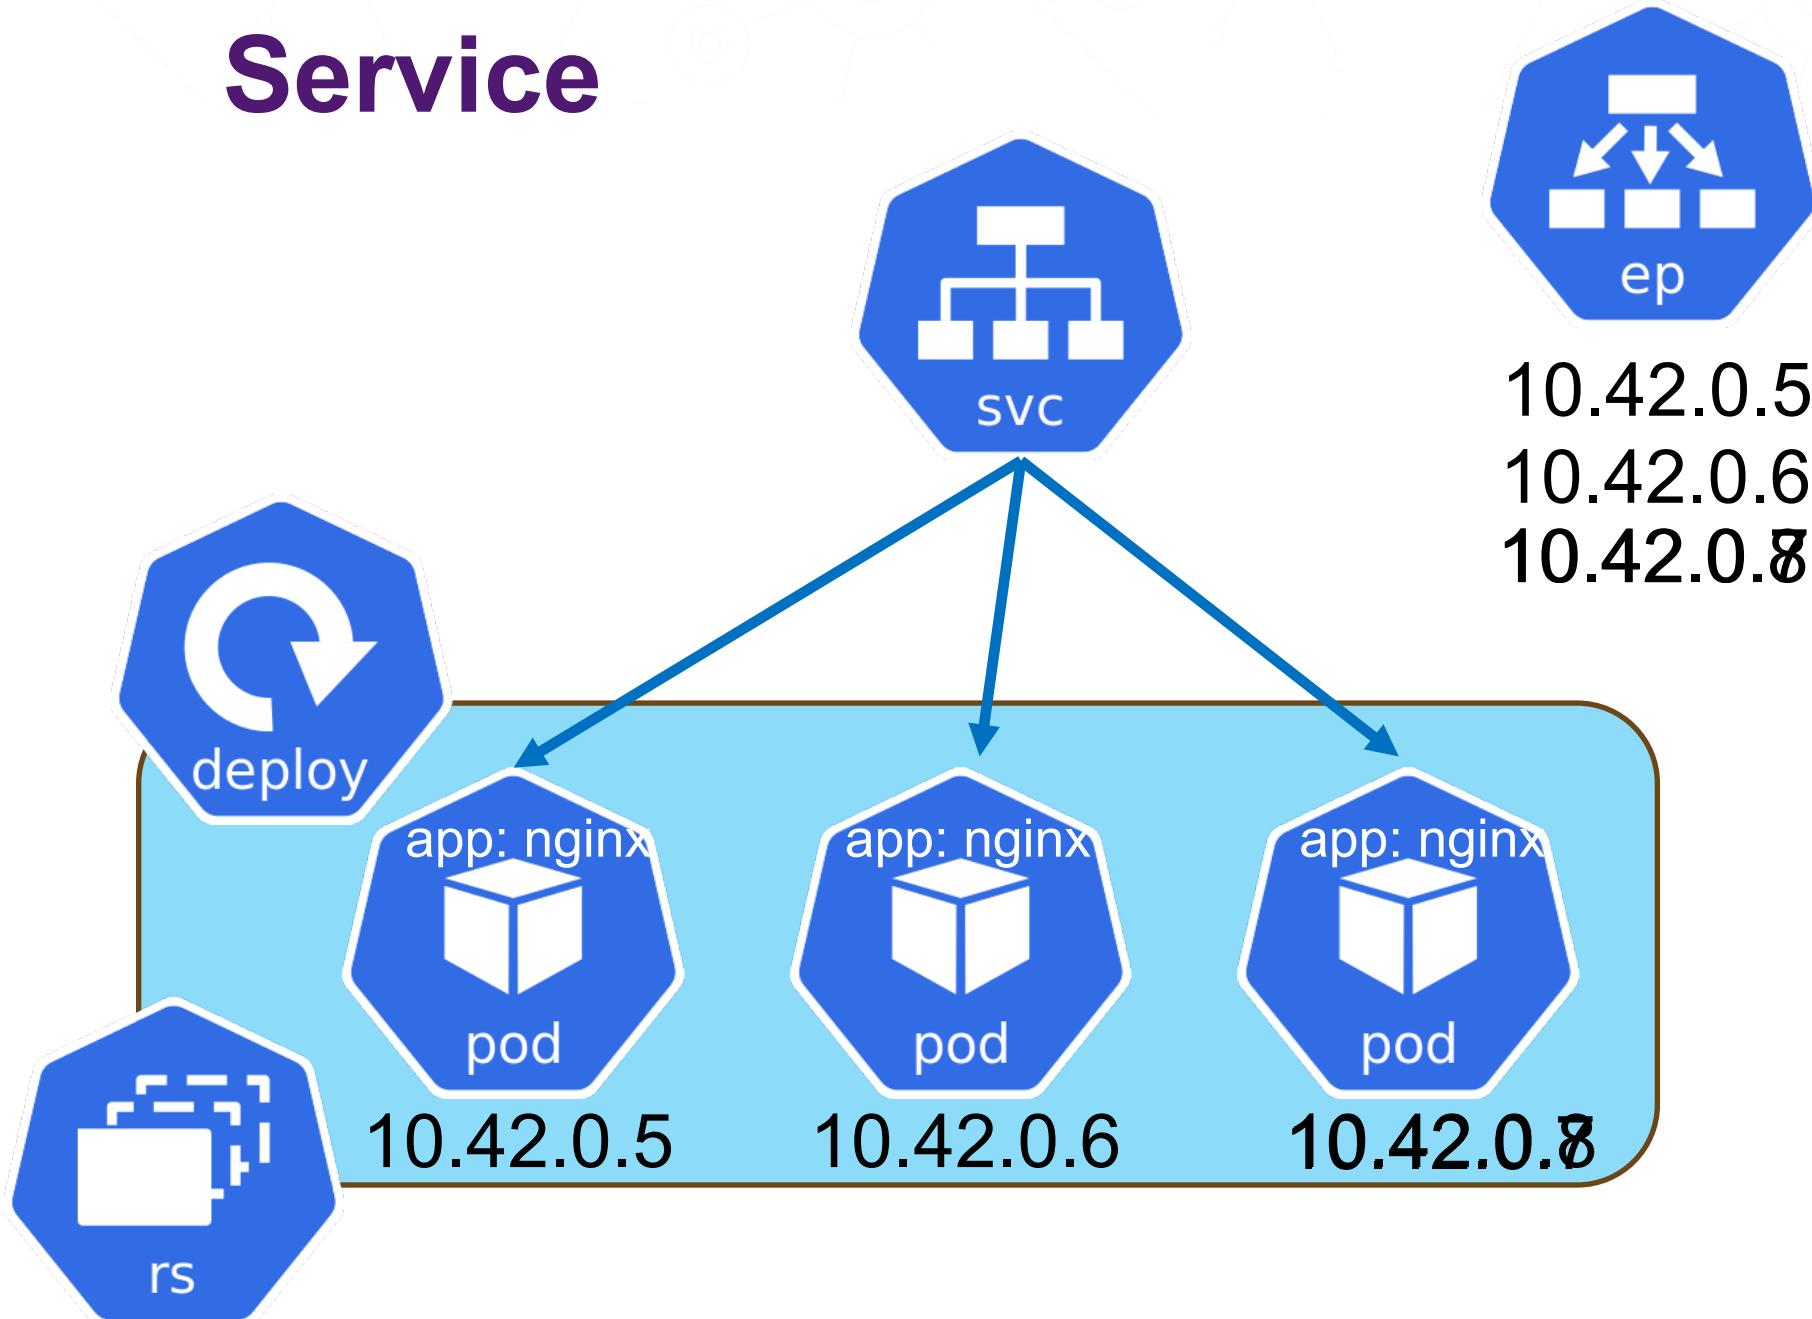
Service



```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
  namespace: mynamespace
spec:
  selector:
    app: nginx
  ports:
  - protocol: TCP
    port: 80
    targetPort: 80
```

- An abstraction to expose a group of Pods
- Service defines a logical set of Endpoints (Pods)
- Pods exposed by a Service are targeted with a label selector
- Types of Services:
 - ClusterIP (inter-cluster, CIDR 10.43.0.0/16)
 - NodePort (port 30000 to 32767)
 - LoadBalancer (cloud provider or external LB)

Service



```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
  namespace: mynamespace
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```



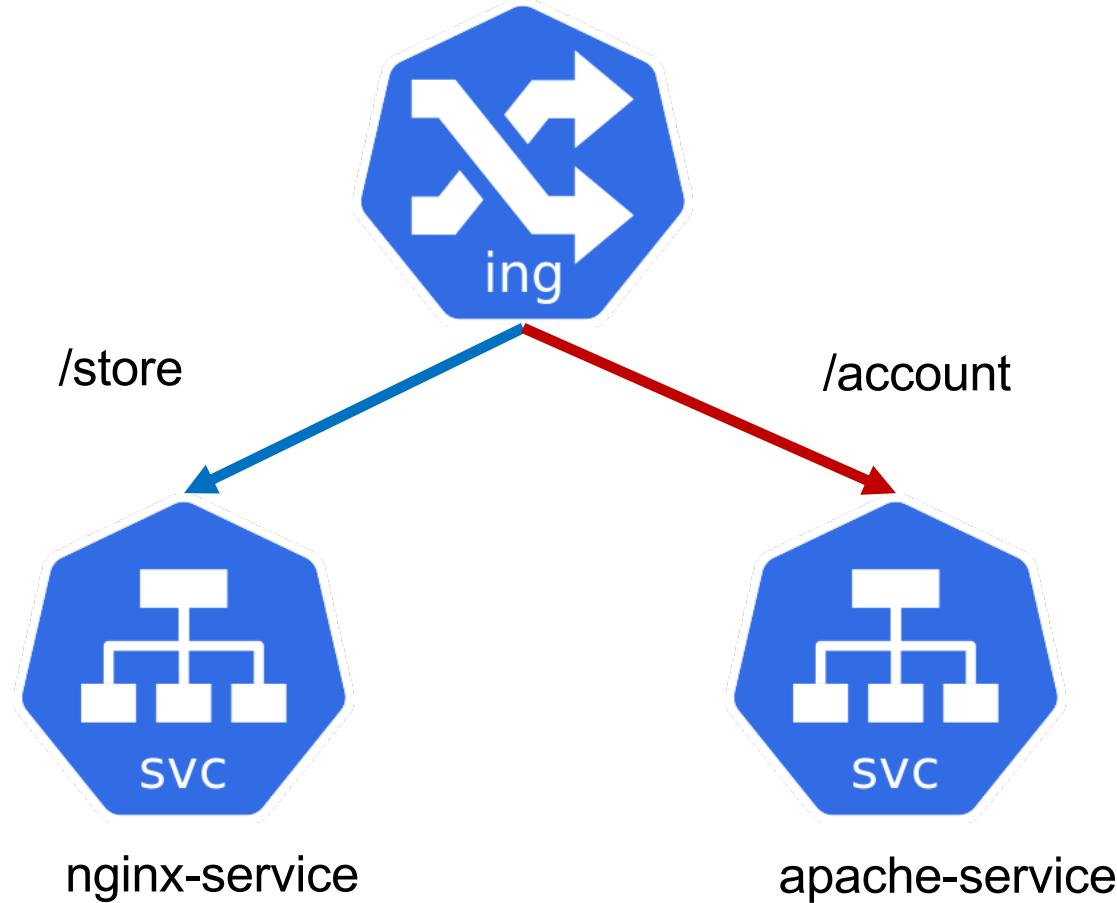
Ingress



- External access to Services over HTTP/HTTPS with an externally reachable URL
- Load balances traffic
- Terminates SSL/TLS

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: myapp-ingress
  namespace: mynamespace
spec:
  rules:
    - host: myapp.com
      http:
        paths:
          - path: /store
            backend:
              serviceName: nginx-service
              servicePort: 80
          - path: /account
            backend:
              serviceName: apache-service
              servicePort: 8080
```

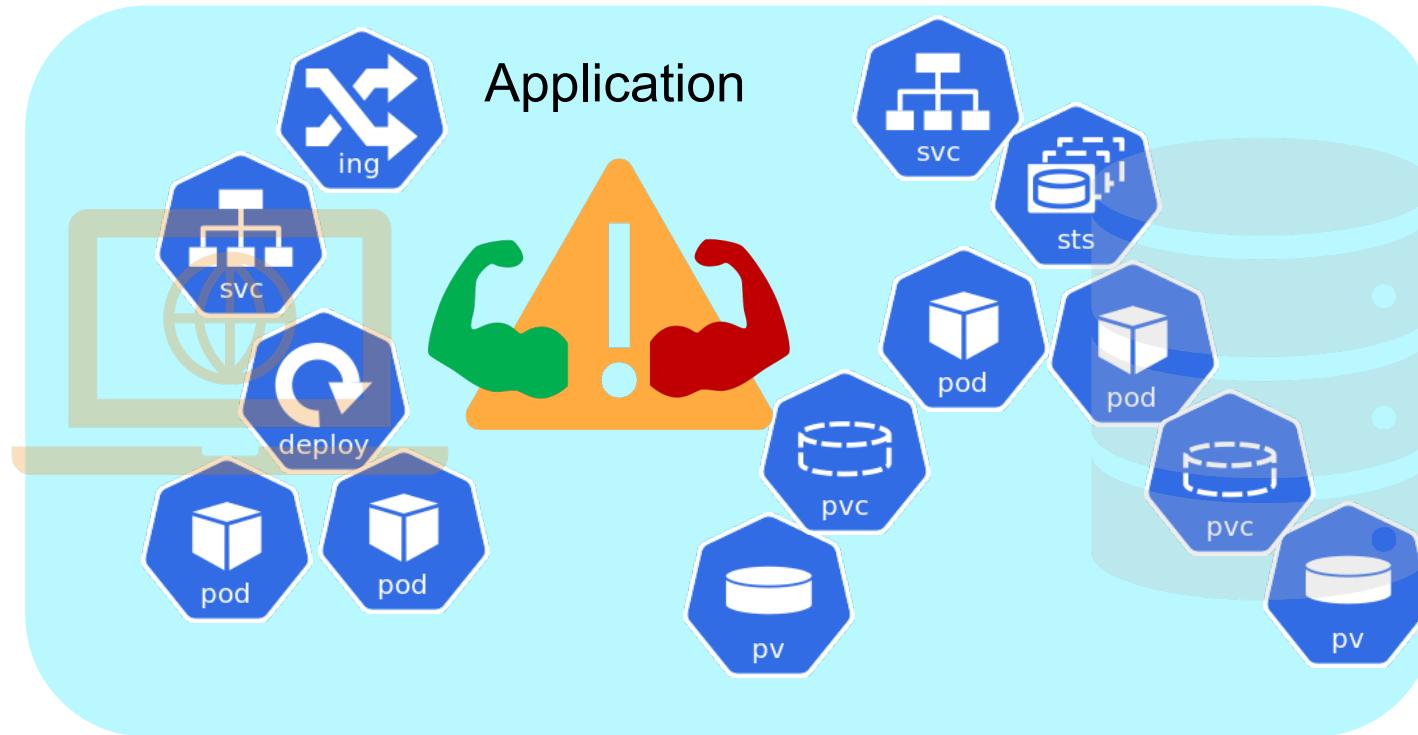
Ingress



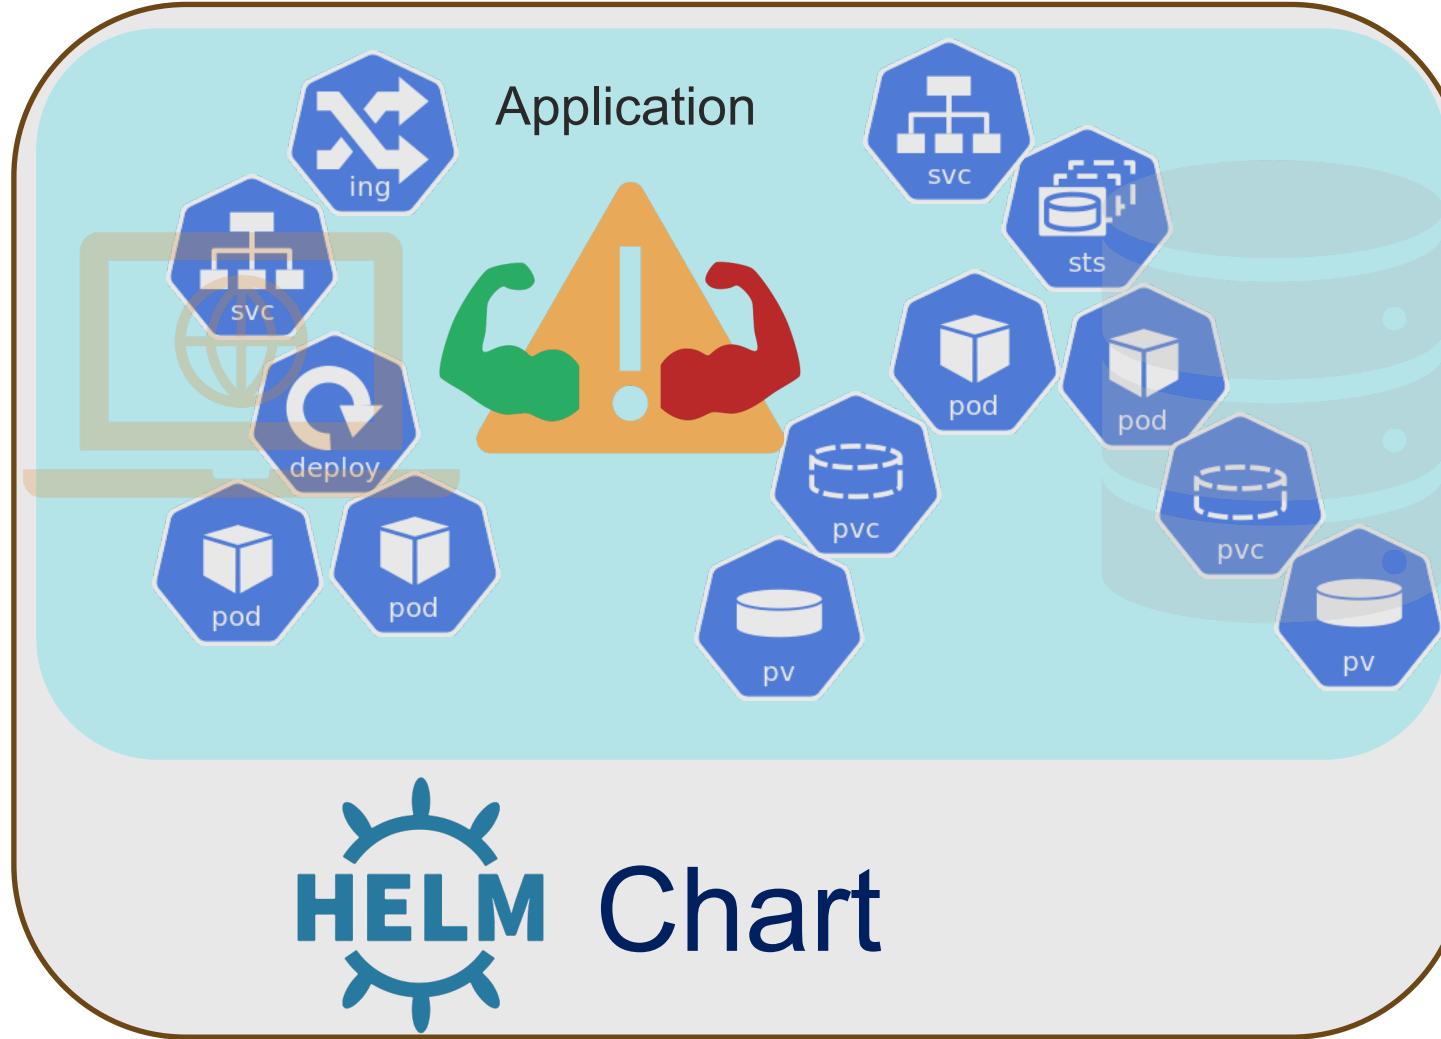
```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: myapp-ingress
  namespace: mynamespace
spec:
  rules:
  - host: myapp.com
    http:
      paths:
      - path: /store
        backend:
          serviceName: nginx-service
          servicePort: 80
      - path: /account
        backend:
          serviceName: apache-service
          servicePort: 8080
```



Application Definition



Application Definition



Test Environments



minikube



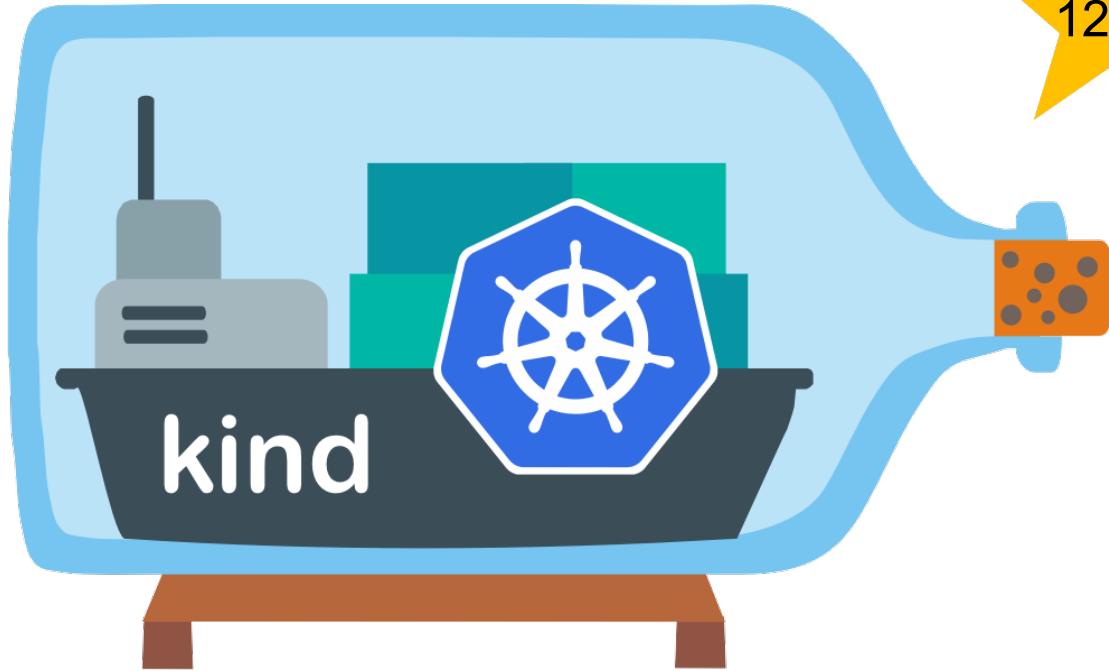
<https://minikube.sigs.k8s.io/docs/start/>

Implements a local cluster on Linux, Windows, macOS

Pre-reqs: container runtime or virtual machine manager

- Podman
- Docker
- QEMU
- Hyperkit
- Hyper-V
- KVM
- Parallels
- VirtualBox
- VMware Fusion/Workstation

kind



- Run local Kubernetes clusters using Docker container “nodes” on MacOS, Windows, Linux
- Used by the Kubernetes project to test and run integration tests

Pre-reqs:

- Docker
- kubectl

<https://kind.sigs.k8s.io/docs/user/quick-start>



<https://k3s.io/>

- CNCF Project
- Great for resource-constrained environments
- Linux only
- Uses an etcd shim called kine (kine is not etcd)
- Can use etcd, PostgreSQL, MySQL, SQLite, MariaDB, NATS JetStream

K3s



- Wrapper around k3s to run in Docker

Pre-reqs:

- Docker
- kubectl

<https://k3d.io/>



Tutorial



Tutorial



<https://github.com/reylejano/kubernetes-essentials-devweek>

- Deploy a cluster
- Deploy containerized application in a Pod
- Deploy the containerized application in a Deployment
- Update the application
- Scale the Deployment up and down
- Get traffic to our application from outside the cluster

<https://github.com/reylejano/kubernetes-essentials-devweek>

Thank you

